

The Feasibility of Implementing Maximum Flow and Minimum-Cost Flow in Almost-Linear Time with an IPM

Exam presentation

Adrian Borup, Albert Rise Nielsen

Research Project
IT University of Copenhagen

December 18, 2024

Preamble

- Chen, Kyng, Liu, Peng, Gutenberg, and Sachdeva presented a breakthrough in 2022
- Solves minimum-cost flow and maximum flow in $m^{1+o(1)}$ time
- Two key components:
 - 1 Interior-point method (IPM) solving $m^{1+o(1)}$ subproblems.
 - 2 A dynamic graph data structure to solve subproblems efficiently.
- No known implementation.
- Is the algorithm feasible in practice?

Table of Contents

- 1 Background: Minimum-cost flow
- 2 Interior-point method
 - Min-cost flow as a linear program
 - An IPM for min-cost flow
 - Min-ratio cycle subproblem
 - Combining everything
- 3 Data structure
- 4 Results and discussion
- 5 Conclusion and scope for masters thesis

Table of Contents

- 1 Background: Minimum-cost flow
- 2 Interior-point method
 - Min-cost flow as a linear program
 - An IPM for min-cost flow
 - Min-ratio cycle subproblem
 - Combining everything
- 3 Data structure
- 4 Results and discussion
- 5 Conclusion and scope for masters thesis

Background: Minimum-cost flow

“the cheapest possible way of sending flow, within some capacities, through a flow network”

Background: Minimum-cost flow

- Directed graph
- Lower and upper capacities
- Cost to use: $c \cdot f$
- Vertex demands
- Minimise cost while respecting capacities and demands




Table of Contents

- 1 Background: Minimum-cost flow
- 2 Interior-point method
 - Min-cost flow as a linear program
 - An IPM for min-cost flow
 - Min-ratio cycle subproblem
 - Combining everything
- 3 Data structure
- 4 Results and discussion
- 5 Conclusion and scope for masters thesis

Min-cost flow as a linear program

Define:

$\mathbf{f} \in \mathbb{R}^E$ the flow in each edge

$\mathbf{c} \in \mathbb{R}^E$ the cost of each edge


$\mathbf{u}_e^-, \mathbf{u}_e^+ \in \mathbb{R}^E$ the lower/upper capacity of each edge

$\mathbf{d} \in \mathbb{R}^V$ the demand of each vertex

$\mathbf{B} \in \mathbb{R}^{E \times V}$ edge-vertex incidence matrix

$$\begin{array}{ll}
 \text{minimize} & \mathbf{c}^\top \mathbf{f} & (\text{minimum cost}) \\
 \text{subject to} & \mathbf{B}^\top \mathbf{f} = \mathbf{d} & (\text{vertex demand constraints}) \\
 & \text{and } \mathbf{u}_e^- \leq \mathbf{f}_e \leq \mathbf{u}_e^+ & (\text{edge capacity constraints})
 \end{array}$$

Optimising a linear program: the interior-point method

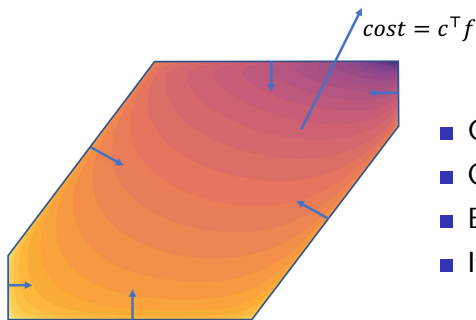


A potential-reduction IPM for min-cost flow

$$\Phi(\mathbf{f}) = \underbrace{m \log(\mathbf{c}^\top \mathbf{f} - F^*)}_{\text{Objective}} - \underbrace{\sum_{e \in E} (\log(\mathbf{u}_e^+ - \mathbf{f}_e) + \log(\mathbf{f}_e - \mathbf{u}_e^-))}_{\text{Barrier}}$$

A potential-reduction IPM for min-cost flow

$$\Phi(\mathbf{f}) = \underbrace{m \log(\mathbf{c}^\top \mathbf{f} - F^*)}_{\text{Objective}} - \underbrace{\sum_{e \in E} (\log(\mathbf{u}_e^+ - \mathbf{f}_e) + \log(\mathbf{f}_e - \mathbf{u}_e^-))}_{\text{Barrier}}$$



- Goal: minimise $\Phi(\mathbf{f})$
- Objective \rightarrow minimises cost
- Barrier \rightarrow avoids boundary
- It's a balance

What do we need?

- 1 An initial point: some feasible flow
- 2 A step function: find some flow that decreases the potential


Dropping the potential

Question: How do we find which direction decreases $\Phi(\mathbf{f})$?

Dropping the potential

Question: How do we find which direction decreases $\Phi(\mathbf{f})$?

Answer: $\nabla\Phi(\mathbf{f})$, the gradient of $\Phi(\mathbf{f})$



Dropping the potential

Question: How do we find a *valid flow* that decreases $\Phi(\mathbf{f})$?

Dropping the potential

Question: How do we find a *valid flow* that decreases $\Phi(\mathbf{f})$?

Idea:

- We have the constraint $\mathbf{B}^\top \mathbf{f} = \mathbf{d}$
- Find a circulation Δ such that $\mathbf{B}^\top \Delta = \mathbf{0}$
Augmenting Δ gives $\mathbf{B}^\top (\mathbf{f} + \Delta) = \mathbf{d}$, still respecting demands
- Goal: find a Δ that steps in the direction of the gradient.

Minimum-ratio cycles (in general)

Define:

C a cycle consisting of a set of edges

$c(e)$ the cost of an edge

$t(e)$ the time of an edge

Minimum-ratio cycles (in general)

Define:

C a cycle consisting of a set of edges

$c(e)$ the cost of an edge

$t(e)$ the time of an edge

Find the cycle that minimises:

$$\sum_{e \in C} c(e) / \sum_{e \in C} t(e)$$

The min-ratio cycle subproblem in the IPM

$\mathbf{g} \in \mathbb{R}^E$ edge gradients, $\mathbf{g} = \nabla \Phi(\mathbf{f})$

$\ell \in \mathbb{R}_{>0}^E$ edge lengths, $\ell_e = 1 / (\min(\mathbf{u}_e^+ - \mathbf{f}_e, \mathbf{f}_e - \mathbf{u}_e^-))$

$\mathbf{L} \in \mathbb{R}^{E \times E}$ $\mathbf{L} = \text{diag}(\ell)$



$$\min_{\mathbf{B}^\top \Delta = 0} \frac{\mathbf{g}^\top \Delta}{\|\mathbf{L} \Delta\|_1}$$

In simpler terms:

Find a circulation Δ that minimises $\frac{\sum_e \mathbf{g}_e \cdot \Delta_e}{\sum_e |\ell_e \cdot \Delta_e|}$

Examples

$$\min_{\mathbf{B}^\top \Delta = 0} \frac{\mathbf{g}^\top \Delta}{\|\mathbf{L}\Delta\|_1}$$





$$\mathbf{g}^\top \Delta = 1 + (-2) + (-3) = -4$$

$$\|\mathbf{L}\Delta\|_1 = 2 + 1 + 2 = 5$$

Examples

$$\min_{\mathbf{B}^\top \Delta = 0} \frac{\mathbf{g}^\top \Delta}{\|\mathbf{L}\Delta\|_1}$$



$$\begin{aligned}\mathbf{g}^\top \Delta &= -1 + 2 + 3 = 4 \\ \|\mathbf{L}\Delta\|_1 &= 2 + 1 + 2 = 5\end{aligned}$$

Why min-ratio cycles?

What are the benefits of using min-ratio cycles?

- 1 Gradients ensure good step directions for IPM
- 2 Outputs simple cycles
- 3 [CKLPGS22] can compute *approximate* min-ratio cycles very efficiently

Combining everything

Given some initial flow $\mathbf{f}^{(0)}$, perform iterations:

- 1 Compute $\mathbf{g}(\mathbf{f}^{(t)})$ and $\ell(\mathbf{f}^{(t)})$

Combining everything

Given some initial flow $\mathbf{f}^{(0)}$, perform iterations:

- 1 Compute $\mathbf{g}(\mathbf{f}^{(t)})$ and $\ell(\mathbf{f}^{(t)})$
- 2 Find min-ratio cycle Δ with $\min_{\mathbf{B}^\top \Delta = 0} \frac{\mathbf{g}^\top \Delta}{\|\mathbf{L}\Delta\|_1}$

Combining everything

Given some initial flow $\mathbf{f}^{(0)}$, perform iterations:

- 1 Compute $\mathbf{g}(\mathbf{f}^{(t)})$ and $\ell(\mathbf{f}^{(t)})$
- 2 Find min-ratio cycle Δ with $\min_{\mathbf{B}^\top \Delta = 0} \frac{\mathbf{g}^\top \Delta}{\|\mathbf{L}\Delta\|_1}$
- 3 Compute scaling factor $\eta = -\kappa^2 / (50 \cdot \mathbf{g}^\top \Delta)$

Combining everything

Given some initial flow $\mathbf{f}^{(0)}$, perform iterations:

- 1 Compute $\mathbf{g}(\mathbf{f}^{(t)})$ and $\ell(\mathbf{f}^{(t)})$
- 2 Find min-ratio cycle Δ with $\min_{\mathbf{B}^\top \Delta = 0} \frac{\mathbf{g}^\top \Delta}{\|\mathbf{L}\Delta\|_1}$
- 3 Compute scaling factor $\eta = -\kappa^2 / (50 \cdot \mathbf{g}^\top \Delta)$
- 4 Update the flow, $\mathbf{f}^{(t+1)} \leftarrow \mathbf{f}^{(t)} + \eta \Delta$

Combining everything

Given some initial flow $\mathbf{f}^{(0)}$, perform iterations:

- 1 Compute $\mathbf{g}(\mathbf{f}^{(t)})$ and $\ell(\mathbf{f}^{(t)})$
- 2 Find min-ratio cycle Δ with $\min_{\mathbf{B}^\top \Delta = 0} \frac{\mathbf{g}^\top \Delta}{\|\mathbf{L}\Delta\|_1}$
- 3 Compute scaling factor $\eta = -\kappa^2 / (50 \cdot \mathbf{g}^\top \Delta)$
- 4 Update the flow, $\mathbf{f}^{(t+1)} \leftarrow \mathbf{f}^{(t)} + \eta \Delta$

Run $m^{1+o(1)}$ iterations, then round the result.

Running time

Problem: exact \mathbf{g} , ℓ , and min-ratio cycles are expensive to compute.

Solution: use *approximations*.

- └ Interior-point method
- └ Combining everything

Running time

Problem: exact \mathbf{g} , ℓ , and min-ratio cycles are expensive to compute.

Solution: use *approximations*.

Claims by [CKLPGS22]:

- 1 IPM can terminate in $m^{1+o(1)}$ iterations
- 2 An $m^{o(1)}$ -approximate solution for min-ratio cycle suffices
- 3 Can compute $m^{o(1)}$ -approximate solution in amortized $m^{o(1)}$ time
- 4 At most $m^{1+o(1)}$ changes to \mathbf{g} and ℓ

Running time

Problem: exact \mathbf{g} , ℓ , and min-ratio cycles are expensive to compute.

Solution: use *approximations*.

Claims by [CKLPGS22]:

- 1 IPM can terminate in $m^{1+o(1)}$ iterations
- 2 An $m^{o(1)}$ -approximate solution for min-ratio cycle suffices
- 3 Can compute $m^{o(1)}$ -approximate solution in amortized $m^{o(1)}$ time
- 4 At most $m^{1+o(1)}$ changes to \mathbf{g} and ℓ

$m^{1+o(1)}$ iterations, each running in amortized $m^{o(1)}$ time.

Total $m^{1+o(1)}$ time.

Extras


- How do we find an initial feasible flow? $O(m)$ time algorithm described in report.
- IPM assumes that optimal cost is known. When unknown, binary search can be applied. Described in report.

Table of Contents

- 1 Background: Minimum-cost flow
- 2 Interior-point method
 - Min-cost flow as a linear program
 - An IPM for min-cost flow
 - Min-ratio cycle subproblem
 - Combining everything
- 3 Data structure
- 4 Results and discussion
- 5 Conclusion and scope for masters thesis

Dynamic min-ratio cycle

The fundamental idea of the data structure is to reduce the number of vertices and edges as much as possible while supporting rebuilds, slowly-changing \mathbf{g} , ℓ , and graph updates.



Dynamic min-ratio cycle

The goal is to approximately solve $\min_{\mathbf{B}^\top \Delta = 0} \frac{\mathbf{g}^\top \Delta}{\|\mathbf{L}\Delta\|_1}$ with the help of a specialized data structure

Algorithm:

- 1 Sample a random “low-stretch spanning tree” T
- 2 Return the best “tree cycle” in T ($\text{cycle}_T(e)$)
- 3 Repeat $O(\log n)$ times




Table of Contents

- 1 Background: Minimum-cost flow
- 2 Interior-point method
 - Min-cost flow as a linear program
 - An IPM for min-cost flow
 - Min-ratio cycle subproblem
 - Combining everything
- 3 Data structure
- 4 Results and discussion
- 5 Conclusion and scope for masters thesis

Implementation

We implement a partial, static algorithm:

- Focus on IPM
- Ignore dynamic data structure

Still allows us to evaluate claims about running time via the number of iterations.

Implementation

We implement a partial, static algorithm:

- Focus on IPM
- Ignore dynamic data structure

Still allows us to evaluate claims about running time via the number of iterations.

Differences from [CKLPGS22]:


- We use an *inefficient, exact* oracle for minimum-ratio cycles.
- We take larger steps than suggested
- We devise our own binary search
- ... see more in report

Correctness

Test suite:

- Focus on maximum flow. Tests from various sources.
- Almost all test cases yield correct maximum flow value.
- Some tests fail flow conservation checks.
- Some crash due to overflow.

Problems: fractional flows



Other problems

- Numeric overflows
- Floating point numbers and precision
- Off-by-one error in binary-search

Experimental results

Grows almost linearly in m , but the constants are huge.




Figure: All graphs are directed acyclic graphs




Figure: All graphs are fully connected graphs

Experimental results

[CKLPGS22] claims running time of $m^{1+o(1)} \log U$ for max flow.

The impact of U on the running time when binary-searching:

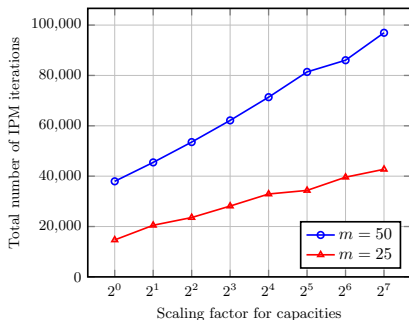


Figure: The total number of IPM iterations for each scaling factor, shown with a logarithmic x-scale.

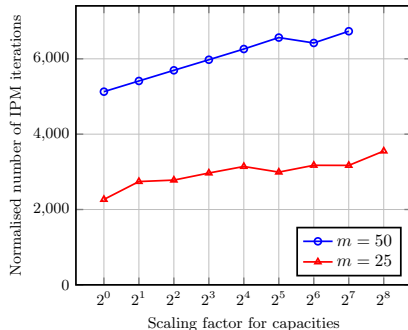


Figure: The number of IPM iterations, normalised by dividing by $\log_2(F_{\max})$ where F_{\max} is the upper limit for the binary search.

Experimental results

Edge updates compared to other flow algorithms.


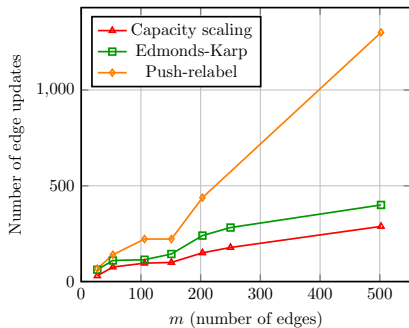


Table of Contents

- 1 Background: Minimum-cost flow
- 2 Interior-point method
 - Min-cost flow as a linear program
 - An IPM for min-cost flow
 - Min-ratio cycle subproblem
 - Combining everything
- 3 Data structure
- 4 Results and discussion
- 5 Conclusion and scope for masters thesis

Conclusion and scope for masters thesis

- We give a high-level overview of the theory behind [CKLPGS22]
- We provide a static algorithm for the IPM
- We evaluate the algorithm on its usefulness in practice

Conclusion and scope for masters thesis

- We give a high-level overview of the theory behind [CKLPGS22]
- We provide a static algorithm for the IPM
- We evaluate the algorithm on its usefulness in practice
- We suggest that [CKLPGS22] is not very practical.
 - Continuous optimisation is problematic in practice
 - High number of iterations
 - Precision issues
 - Complex
- Scope for master's thesis
 - We could attempt to implement the data structure.
 - Work by Bernstein, Blikstad, Saranurak, and Tu ([BBST24]) seems promising and implementable.

Questions

It's time for questions.

References

- [BBST24] Aaron Bernstein, Joakim Blikstad, Thatchaphol Saranurak, and Ta-Wei Tu. *Maximum Flow by Augmenting Paths in $\tilde{O}(n^{2+o(1)})$ Time*. arXiv:2406.03648. June 2024. DOI: 10.48550/arXiv.2406.03648. URL: <http://arxiv.org/abs/2406.03648> (visited on 11/13/2024).
- [CKLPGS22] Li Chen, Rasmus Kyng, Yang P. Liu, Richard Peng, Maximilian Probst Gutenberg, and Sushant Sachdeva. *Maximum Flow and Minimum-Cost Flow in Almost-Linear Time*. arXiv:2203.00671. Apr. 2022. DOI: 10.48550/arXiv.2203.00671. URL: <http://arxiv.org/abs/2203.00671> (visited on 11/13/2024).
- [Kar84] Narendra Karmarkar. “A New Polynomial-Time Algorithm for Linear Programming-II”. In: *Combinatorica* 4 (Dec. 1984), pp. 302–311. DOI: 10.1007/BF02579150.
- [LC22] Yang P. Liu and Li Chen. *Maximum Flow and Minimum-Cost Flow in Almost Linear Time*. Stanford, Mar. 2022. URL: https://yangpliu.github.io/pdf/Flows_Stanford.pdf.